

Predicting film box office openings with Wikipedia

Andrew Lim

14 Jan 2013

Abstract

To what extent can Wikipedia be used to measure social interest? In this analysis, I examine the ability of Wikipedia edit data to predict opening box office revenues for US films. I present a simple model fit to films released during 2007-2011 based on features in their Wikipedia articles. While this model’s predictive power is probably insufficient for all but the roughest estimates of revenue, it does demonstrate how popular interest in films is reflected in Wikipedia activity.

1 Introduction: Wikipedia as a gauge of social interest




Figure 1: Count of Wikipedia article edits for the films used in this paper’s training dataset over the 4 weeks prior each film’s respective release date, bucketed by days before release date that the edits occurred. This graph shows the uptick in editing activity that typically accompanies a film’s release.

According to its article about itself (as of this writing), Wikipedia is “a collaboratively edited, multi-lingual, free Internet encyclopedia” launched in January 2001. [6] Its articles can be edited by anyone, either anonymously (though the editor’s IP address is logged) or with a registered user account. The edit history of each article is saved with a timestamp. Interested users can view any past version of

an article, and an article’s edit history exhibits an evolving record of Wikipedia’s “knowledge”¹ of its subject.

As such, Wikipedia’s edit history can be viewed as a barometer of social interest. For example, when a person is in the news, edit activity on his or her article often spikes. In fact, Wikipedia has template warnings indicating when an article is likely to be in flux due to a relevant current event. Edit activity on Wikipedia, in this sense, is akin to mentions on social networks like Facebook or Twitter, although perhaps with a smaller participating audience (although many people read Wikipedia, not nearly so many participate in its creation).

One area where we can try to gauge the degree to which Wikipedia activity reflects social interest is in film box office performance. Films have relatively well-defined release dates prior to which we can measure activity on Wikipedia. They also have well-defined, measurable outcome - revenues at the ticket booth - that is clearly sensitive to popular interest. Theater owners obviously have a direct financial interest in knowing how well a film is going to perform. Advertisers and publicists, sellers of tie-in products, and film journalists have a slightly more indirect but still strong interest; they will want to know how they should spend their time and money. Can we use Wikipedia to usefully predict films’ opening box office performances?

2 Formulation of problem and data sources

The specific question I set out to answer was: how accurately, using Wikipedia’s help, can we predict the domestic per-theater box office gross of a film released widely in the US, over the first three days² of its release?

The data sources I used to answer this question were:

- Box Office Mojo (<http://www.boxofficemojo.com/>) - contains detailed box office data. I used it to select the universe of films to analyze and as my source for theatrical release dates, number of opening theaters, and revenues. There is no API - I scraped the data with the Python package Beautiful Soup.
- Rotten Tomatoes (<http://www.rottentomatoes.com/>) - a popular movie review aggregator. I used it to obtain descriptive information about films: genres, runtime, MPAA rating, cast and directors, and so on. It offers an API if you register for a key (which is free as of the present writing).
- Wikipedia (English-language) (<http://en.wikipedia.org/>) - MediaWiki, the name of the web application upon which Wikipedia is based, offers an API, no registration or key necessary.

¹Of course, Wikipedia’s highly open policy means both that it contains a stunning breadth of information from contributors with wide-ranging expertise and that said information is sometimes unreliable. For an example that was in the news not long before this paper was written, see [5], or for Wikipedia’s own list of Wikipedia hoaxes, see [4].

²Films traditionally open on Friday, and their “opening” often refers to their gross over the first Friday, Saturday, and Sunday that they are playing. However there are plenty of non-Friday openings as well. Consequently, I’ve stated the problem in terms of the first three days’ worth of grosses.

Much of the work involved in data retrieval and formatting was to ensure that data retrieved from these three sources corresponded to the same film; data from Rotten Tomatoes and Wikipedia was obtained by using their APIs' search functionalities, which can lead to incorrect hits if you are not careful. For example, we want to make sure that Rotten Tomatoes data for the 2012 film "The Lucky One" is not mapped to the 2008 film "The Lucky Ones," or that for the 2010 film "Salt" we do not examine the Wikipedia article for salt, the mineral.³

The universe of films that I considered were those listed on Box Office Mojo as having opened in at least 1000 theaters. I manually excluded a handful of films that were re-releases or limited-engagement special features. I trained my algorithms on films released between 2007 and 2011, inclusive. In total, 689 films were in the training dataset. Data from films as far back as 2002 were used for some of the feature calculations; see the next section for more details. I tested my algorithm on films released in 2012, of which there were 124.

3 Features

3.1 Descriptive features

Descriptive features considered were: year of release, runtime, MPAA rating, whether the film was released on a Friday or not, and membership in genres as defined by Rotten Tomatoes. Rotten Tomatoes has 18 genre labels, listed below. A film can belong to any number of these genres.

- | | | |
|-----------------------------|-----------------------------|-----------------------------|
| • Action & Adventure | • Documentary | • Romance |
| • Animation | • Drama | • Science Fiction & Fantasy |
| • Art House & International | • Horror | • Special Interest |
| • Classics | • Kids & Family | • Sports & Fitness |
| • Comedy | • Musical & Performing Arts | • Television |
| • Cult Movies | • Mystery & Suspense | • Western |

3.2 Wikipedia-based features

For each Wikipedia article, I measured the number of "edit runs" that had occurred during the period 0 to 7 days prior to midnight on the day of the film's release, and also during the period 7 to 28 days prior. I defined an edit run as a sequence of consecutive edits from the same author (identified by IP address if anonymous). Sometimes on Wikipedia the same author commits several edits in a row, presumably as part of a single effort to edit the page, which I wanted to correspondingly treat as a single edit. I generally found this to be a slight improvement over raw edit count in terms of predictive power.

³Box Office Mojo data had to be scraped from HTML, but the HTML was regular and consistent. Rotten Tomatoes has a nice JSON-based API for data retrieval, but its ranking of returns is quirky, sometimes retrieving obscure films or films with similar names (example: Oliver Stone's 2008 biopic "W." was unfindable through search query, even through the website's front end; I had to go to Stone's Rotten Tomatoes page just to find the relevant web page). Wikipedia both has a nice API and solid/consistent lookup, which is all the more impressive given that it contains articles on anything, not just films.

I also extracted a few features from the content of the article revisions themselves. One feature I used was the average size, in bytes, of revisions in the 28-day window. Other features were obtained by scanning the text of the revisions for certain textual patterns. One was a count of the number of article section headings, another was a count of the number of external file references (typically an image or sound file inserted into the article), and the last was a case-insensitive search for the word “IMAX”.

3.3 Revenues of similar films

A natural approach towards predicting the box office performance of a film is to take a look at comparable films; in particular, the natural benchmark for a sequel is its predecessor. To this end, I created a feature consisting of revenues of “similar” films released in the five years preceding each film’s release (hence data as far back as 2002 was involved, even though the training dataset extended only as far back as 2007). The five-year window was arbitrary but I think forms a reasonable bound when comparing for expected box office performance.

Similarity between two films was defined as the geometric mean of the Jaccard⁴ similarity measures of the films’ 1) Rotten Tomatoes genre information and 2) Rotten Tomatoes cast/director information. The Rotten Tomatoes API only returns the first few starring members of each film’s cast, so the metric is not distorted by differing cast sizes. Directors were always treated as a single person, even if there were co-directors, so for our purposes the Coen brothers, for example, count as a single person.

Iron Man 2	0.4472
Fantastic Four: Rise of the Silver Surfer	0.2462
Iron Man	0.2462
Thor	0.2462
Captain America: The First Avenger	0.1741
Push	0.1741
Sherlock Holmes: A Game of Shadows	0.1741
The Losers	0.1508
Scott Pilgrim vs. the World	0.1508
Sherlock Holmes	0.1508
Shutter Island	0.1508

Figure 2: Example similarity scores: for “The Avengers” (2012).

The feature incorporated into the algorithms was, for each film, the opening revenue of all other films in our universe released up to five prior to that film, weighted by similarity. See Figure 2 for an example of similarity scores for one of the films in the test dataset.

4 Analysis and prediction

I tried a few different prediction algorithms; the one that proved the most effective on the test set, as measured by R^2 , was gradient boosting trees.⁵ Gradient boosting is a general predictive technique pioneered by Jerome Friedman of Stanford in which a predictive formula is generated by summing so-called “weak predictors” that are sequentially fit to the gradient of a specified loss function (for example,

⁴The Jaccard similarity of two sets A and B is defined as $|A \cap B|/|A \cup B|$.

⁵Random forests and ordinary linear regression performed worse, but not by much. Despite the clearly non-normal distribution of the revenue per theater (it has a positive skew), I did not have better success with a generalized linear regression than with ordinary linear regression.

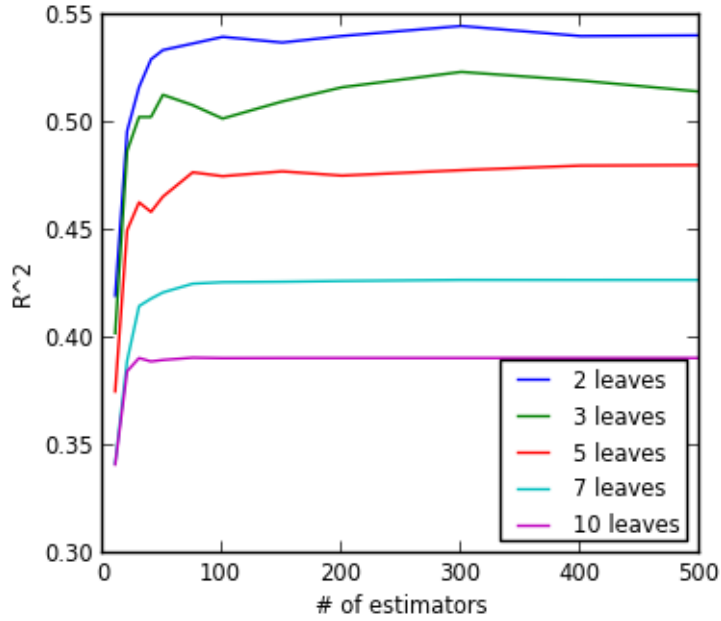


Figure 3: R^2 of gradient boosting tree models on the test dataset as a function of the number of estimator iterations. The different curves represent different numbers of leaves in the weak learner decision trees. The simplest weak learner, a 2-leaf tree, performs the best. Using stochastic gradient boosting trees, in which a subsample of the features are used to fit the decision trees, improved the high-leaf models to some degree. This suggests that the inferior performance of the higher-leaf models may be due to overfitting.

squared error). The overall model may be accurate and robust even if each individual weak predictor is very simplistic. Gradient boosting trees refers to gradient boosting with decision trees as our weak predictors. For details, see Friedman’s article [2], and also Wikipedia’s own page on gradient boosting [3].

I used the Python statistical package `scikit-learn`’s implementation of gradient boosting trees, using the default learning rate and least squares as my loss function. There are a few other model parameters that can be controlled in the user; the most important ones are the number of estimators (the number of weak predictors to fit) and the depth of the trees (how many leaves are in each decision tree - this parameterizes the complexity of each individual weak predictor).

Adapting the example in `scikit-learn`’s documentation [1], I calculated the R^2 of gradient boosting trees at different iterations and tree depths. I fit the model at different parameterizations to the test data. Figure 3 illustrates the results, and shows that this model fits the test data best with about 100 iterations (this is in fact `scikit-learn`’s default value) and a very simple 2-leaf functional form for its weak predictors.

Using a gradient boosting tree model with 100 estimators and 2 leaves in each weak learner, and training on films from 2007 to 2011 as mentioned previously, I was able to achieve an R^2 of 0.5400 on the

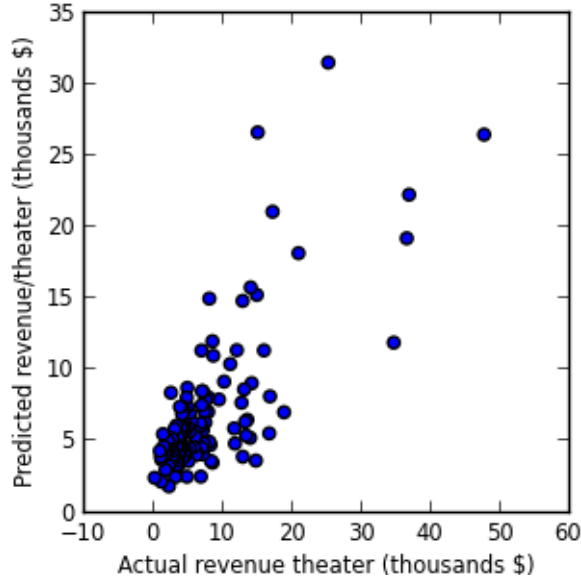


Figure 4: Predicted values vs. actual values.

2012 dataset. The predictions and results are listed in an appendix at the end of this paper. Figure 4 shows a scatter of predictions and actual values.

The frequency with which a feature is used in the model’s decision trees is representative of its importance in generating predictions; highly relevant features will be frequently involved in trees and irrelevant features will be involved rarely or not at all. Figure 5 shows the top 10 features. Several features had frequencies of 0, in particular the boolean variables for several of the genre categories, indicating that they could have been completely omitted without impacting the outcomes of this model.

Feature	Frequency (%)
Wikipedia edit runs 7-28 days prior	18.31
Film runtime	14.60
Opening per theater revenue of similar films	13.30
Wikipedia frequency of headers/subheaders	12.07
Wikipedia edit runs 0-7 days prior	10.97
Wikipedia average size of revisions	9.73
Wikipedia frequency of word “IMAX”	5.07
Wikipedia frequency of external files	4.62
Is comedy	3.74
MPAA rating is PG-13	3.17

Figure 5: Top 10 features in the gradient boosting tree model.

The importance of the Wikipedia data in this model can also be seen by removing the Wikipedia features and rerunning the model, which produces a considerable lower R^2 of 0.3434.

5 Conclusion and avenues for further exploration

While the results above do show that Wikipedia activity has some ability to predict box office returns, I do not think the model in this paper is precise enough to be used as anything but a very rough forecasting tool. Wikipedia is just one possible source of data for quantifying social interest; social networks such as Twitter or Facebook are another; frequency of appearance in news headlines is another.⁶ There are many conceivable metrics to gauge popular interest in seeing a film, and a comprehensive model would include data from many sources

In particular, a many-source approach will help overcome the biases that any one source would have. Although Wikipedia is widely known and read and edited by a wide variety of people, it will still be biased to whatever extent that Wikipedia editors do not reflect the population of people who go to the movies. It is my opinion that the best way to improve this model would be to obtain more measurements of popular interest, particularly data sources whose audiences overlap little with Wikipedia editors - measurements of interest among moviegoing demographics that use the Internet relatively infrequently, for example.

Nevertheless, the partial success in predicting box office revenues with Wikipedia demonstrates that it is one potential source of data to consider when gauging interest - and not just in films, but anywhere popular interest is a concern. Wikipedia could be conceivably used as an input for predictions related to interest in news and current events, ticket sales for events other than films, investor sentiments, and many other areas.

6 References and resources

6.1 APIs and software packages

All data work was done in Python. The software packages that I used were:

- NumPy (<http://www.numpy.org/>) - Popular package for doing general mathematical work.
- pandas (<http://pandas.pydata.org/>) - Offers data structures and functions specific to statistical work.
- scikit-learn (<http://scikit-learn.org/>) - Offers a variety of algorithms and useful functions for machine learning, prediction, error checking and validation, and so on.
- BeautifulSoup (<http://www.crummy.com/software/BeautifulSoup/>) - Library to parse HTML.
- matplotlib (<http://matplotlib.org/>) - Used to generate the graphics in this paper.
- simplejson (<http://pypi.python.org/pypi/simplejson/>) - Used to parse the JSON-formatted data returned by the data sources' APIs.

⁶In fact, I found that the number of opening theaters itself has significant predictive power on per-theater revenue. I omitted it mainly because I wanted to specifically examine Wikipedia's ability to measure social interest.

- PyYAML (<http://pyyaml.org/>) - A package to read YAML (Yet Another Markup Language, in my opinion a bit more flexible and human-editor-friendly than JSON). I formatted my scripts' configuration files in YAML.
- Joblib (<http://packages.python.org/joblib/>) - I used this just for its pickling functionality, to save Wikipedia revision data locally so that I wouldn't have to requery the API repeatedly.

Two of the data sources I used, Rotten Tomatoes and Wikipedia, offered APIs. Documentation on those can be found here:

- Rotten Tomatoes API: <http://developer.rottentomatoes.com/>
- Wikipedia API: <http://en.wikipedia.org/w/api.php> - This is the documentation for MediaWiki-based websites that is returned when you query the API without any parameters
- MediaWiki API documentation: (http://www.mediawiki.org/wiki/API:Main_page) - Another handy resource, written more informally.

6.2 Bibliography

- [1] "Ensemble methods." Retrieved 13 Jan 2012. <http://scikit-learn.org/stable/modules/ensemble.html>
- [2] Friedman, Jerome H. (19 Apr 2001). "Greedy Function Approximation: A Gradient Boosting Machine." Retrieved 10 Jan 2012. <http://www-stat.stanford.edu/~jhf/ftp/trebst.pdf>
- [3] "Gradient boosting." Retrieved 13 Jan 2012. http://en.wikipedia.org/wiki/Gradient_boosting
- [4] "List of hoaxes on Wikipedia." Retrieved 10 Jan 2012. http://en.wikipedia.org/wiki/Wikipedia:List_of_hoaxes_on_Wikipedia
- [5] Pfeiffer, Eric (4 Jan 2013). "War is over: Imaginary 'Bicholm' conflict removed from Wikipedia after five years." Retrieved 10 Jan 2012. <http://news.yahoo.com/blogs/sideshow/war-over-imaginary-bicholim-conflict-page-removed-wikipedia-234717353.html>
- [6] "Wikipedia." Retrieved 10 Jan 2012. <http://en.wikipedia.org/wiki/Wikipedia>

Title	Actual	Predicted	Error (actual - predicted)
Marvel's The Avengers	47698	26452	21247
The Hunger Games	36871	22247	14624
The Dark Knight Rises	36532	19194	17338
The Twilight Saga: Breaking Dawn Part 2	34660	11890	22770
Skyfall	25211	31496	-6285
The Hobbit: An Unexpected Journey	20919	18152	2767
Dr. Seuss' The Lorax	18830	7018	11812
The Amazing Spider-Man	17176	21054	-3877
Ted	16800	8127	8673
Think Like a Man	16693	5536	11157
Brave	15928	11344	4584
Prometheus	15032	26608	-11576
Snow White and the Huntsman	14900	15231	-331
The Devil Inside	14763	3633	11129
Madagascar 3: Europe's Most Wanted	14166	9051	5114
Les Misérables	14015	15750	-1735
The Vow	13929	5224	8705
Taken 2	13525	6490	7035
Magic Mike	13354	5383	7971
Flight	13217	6345	6872
Wreck-It Ralph	13070	8619	4451
Safe House	12880	3902	8978
Men in Black 3	12851	14814	-1962
Hotel Transylvania	12697	7685	5012
Ice Age: Continental Drift	12015	11360	654
Madea's Witness Protection	11749	4824	6925
21 Jump Street	11632	5887	5745
Django Unchained	11070	10390	680
The Bourne Legacy	10185	9162	1023
Wrath of the Titans	9438	7910	1528
The Expendables 2	8622	10977	-2355
Contraband	8505	3486	5019
Paranormal Activity 4	8501	11988	-3487
The Campaign	8296	3597	4699
Underworld Awakening	8222	4731	3491
Act of Valor	8054	4942	3112
John Carter	8050	14971	-6920
Dark Shadows	7906	8058	-153
Journey 2: The Mysterious Island	7878	7045	833
Chronicle	7569	7876	-306
Red Tails	7477	6304	1173
The Woman in Black	7311	7075	237
Tyler Perry's Good Deeds	7310	4824	2485
The Lucky One	7137	4066	3072
Sinister	7126	5747	1380
Total Recall	7103	8488	-1385
Resident Evil: Retribution	6989	8491	-1501
Ghost Rider: Spirit of Vengeance	6968	4521	2446
Looper	6952	7526	-573
Battleship	6920	11337	-4417
Project X	6891	6061	830
Chimpanzee	6829	2512	4317
American Reunion	6740	6247	493
The Possession	6297	4131	2166
The Grey	6174	4020	2154
Savages	6095	5649	445
Argo	6020	5313	708
Life of Pi	5805	4569	1236
This Means War	5458	6400	-942

Figure 6: 2012 predictions and errors, sorted by actual revenue per theater

Title	Actual	Predicted	Error (actual - predicted)
Abraham Lincoln: Vampire Hunter	5247	5668	-421
The Cabin in the Woods	5245	6979	-1734
Sparkle	5189	4511	677
Mirror Mirror	5032	3589	1444
Red Dawn	4916	7430	-2514
The Three Stooges	4892	5981	-1089
Rise of the Guardians	4869	8725	-3856
End of Watch	4818	2503	2315
Cloud Atlas	4787	8046	-3259
Step Up Revolution	4570	4409	162
Jack Reacher	4538	5726	-1188
Alex Cross	4489	3955	533
That's My Boy	4440	6258	-1818
Parental Guidance	4392	4140	252
Diary of a Wimpy Kid: Dog Days	4312	5826	-1514
The Dictator	4245	7210	-2965
The Secret World of Arrietty	4235	6930	-2695
The Man with the Iron Fists	4235	6053	-1818
One For the Money	4207	4619	-411
Rock of Ages	4161	7405	-3244
ParaNorman	4108	6899	-2791
Joyful Noise	4104	3791	313
The Watch	4025	5435	-1411
House at the End of The Street	3985	5928	-1942
This Is 40	3976	5154	-1178
Here Comes the Boom	3921	4479	-559
Hope Springs	3850	5357	-1507
Frankenweenie	3798	7396	-3599
Trouble with the Curve	3786	4652	-866
Big Miracle	3645	4006	-361
The Five-Year Engagement	3614	4890	-1276
What to Expect When You're Expecting	3491	5260	-1769
Safe	3483	3176	307
Haywire	3454	3687	-232
The Pirates! Band of Misfits	3317	6083	-2766
The Raven	3309	3317	-8
Chernobyl Diaries	3270	4665	-1395
A Thousand Words	3268	4619	-1351
Wanderlust	3260	3647	-387
Silent House	3136	2481	655
Katy Perry: Part of Me	3031	5843	-2812
The Odd Life of Timothy Green	2954	3826	-872
Seven Psychopaths	2821	4024	-1203
Killing Them Softly	2811	5026	-2215
Silent Hill: Revelation 3D	2735	5153	-2417
Lockout	2700	4342	-1642
Premium Rush	2674	4297	-1623
Man on a Ledge	2669	5169	-2500
Dredd	2505	8384	-5879
Seeking a Friend for the End of the World	2352	4462	-2109
The Collection	2213	1832	381
Gone	2182	3289	-1106
People Like Us	2071	5024	-2953
Playing for Keeps	2027	4417	-2390
Atlas Shrugged: Part II	1731	2991	-1260
Lawless	1700	4400	-2701
The Words	1696	4066	-2370
The Guilt Trip	1448	4619	-3171
Fun Size	1361	5487	-4127
The Cold Light of Day	1212	3987	-2775
Chasing Mavericks	1133	3702	-2569
Last Ounce of Courage	1127	2156	-1029
Won't Back Down	1035	3902	-2867
Hit and Run	910	4311	-3401
Oogieloves In The BIG Balloon Adventure	162	2429	-2267

Figure 7: 2012 predictions and errors (continued), sorted by actual revenue per theater